

Validation loss vs $\log_{10}(\text{training step})$ (classification, full-dataset runs)

